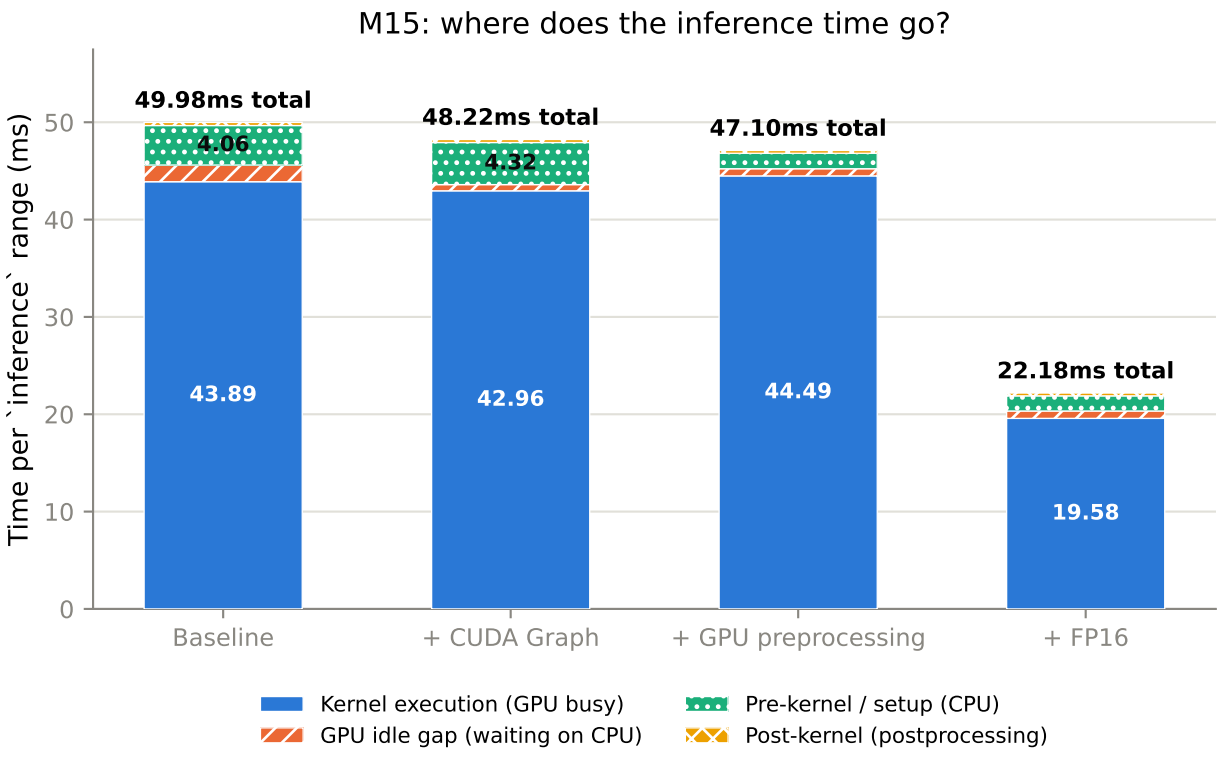
M15: where does the inference time go?
Time per `inference` range (ms)
0
10
20
30
40
50
49.98ms total
4.06
43.89
48.22ms total
4.32
42.96
47.10ms total
44.49
22.18ms total
19.58
Baseline
+ CUDA Graph
+ GPU preprocessing
+ FP16
Kernel execution (GPU busy)
GPU idle gap (waiting on CPU)
Pre-kernel / setup (CPU)
Post-kernel (postprocessing)