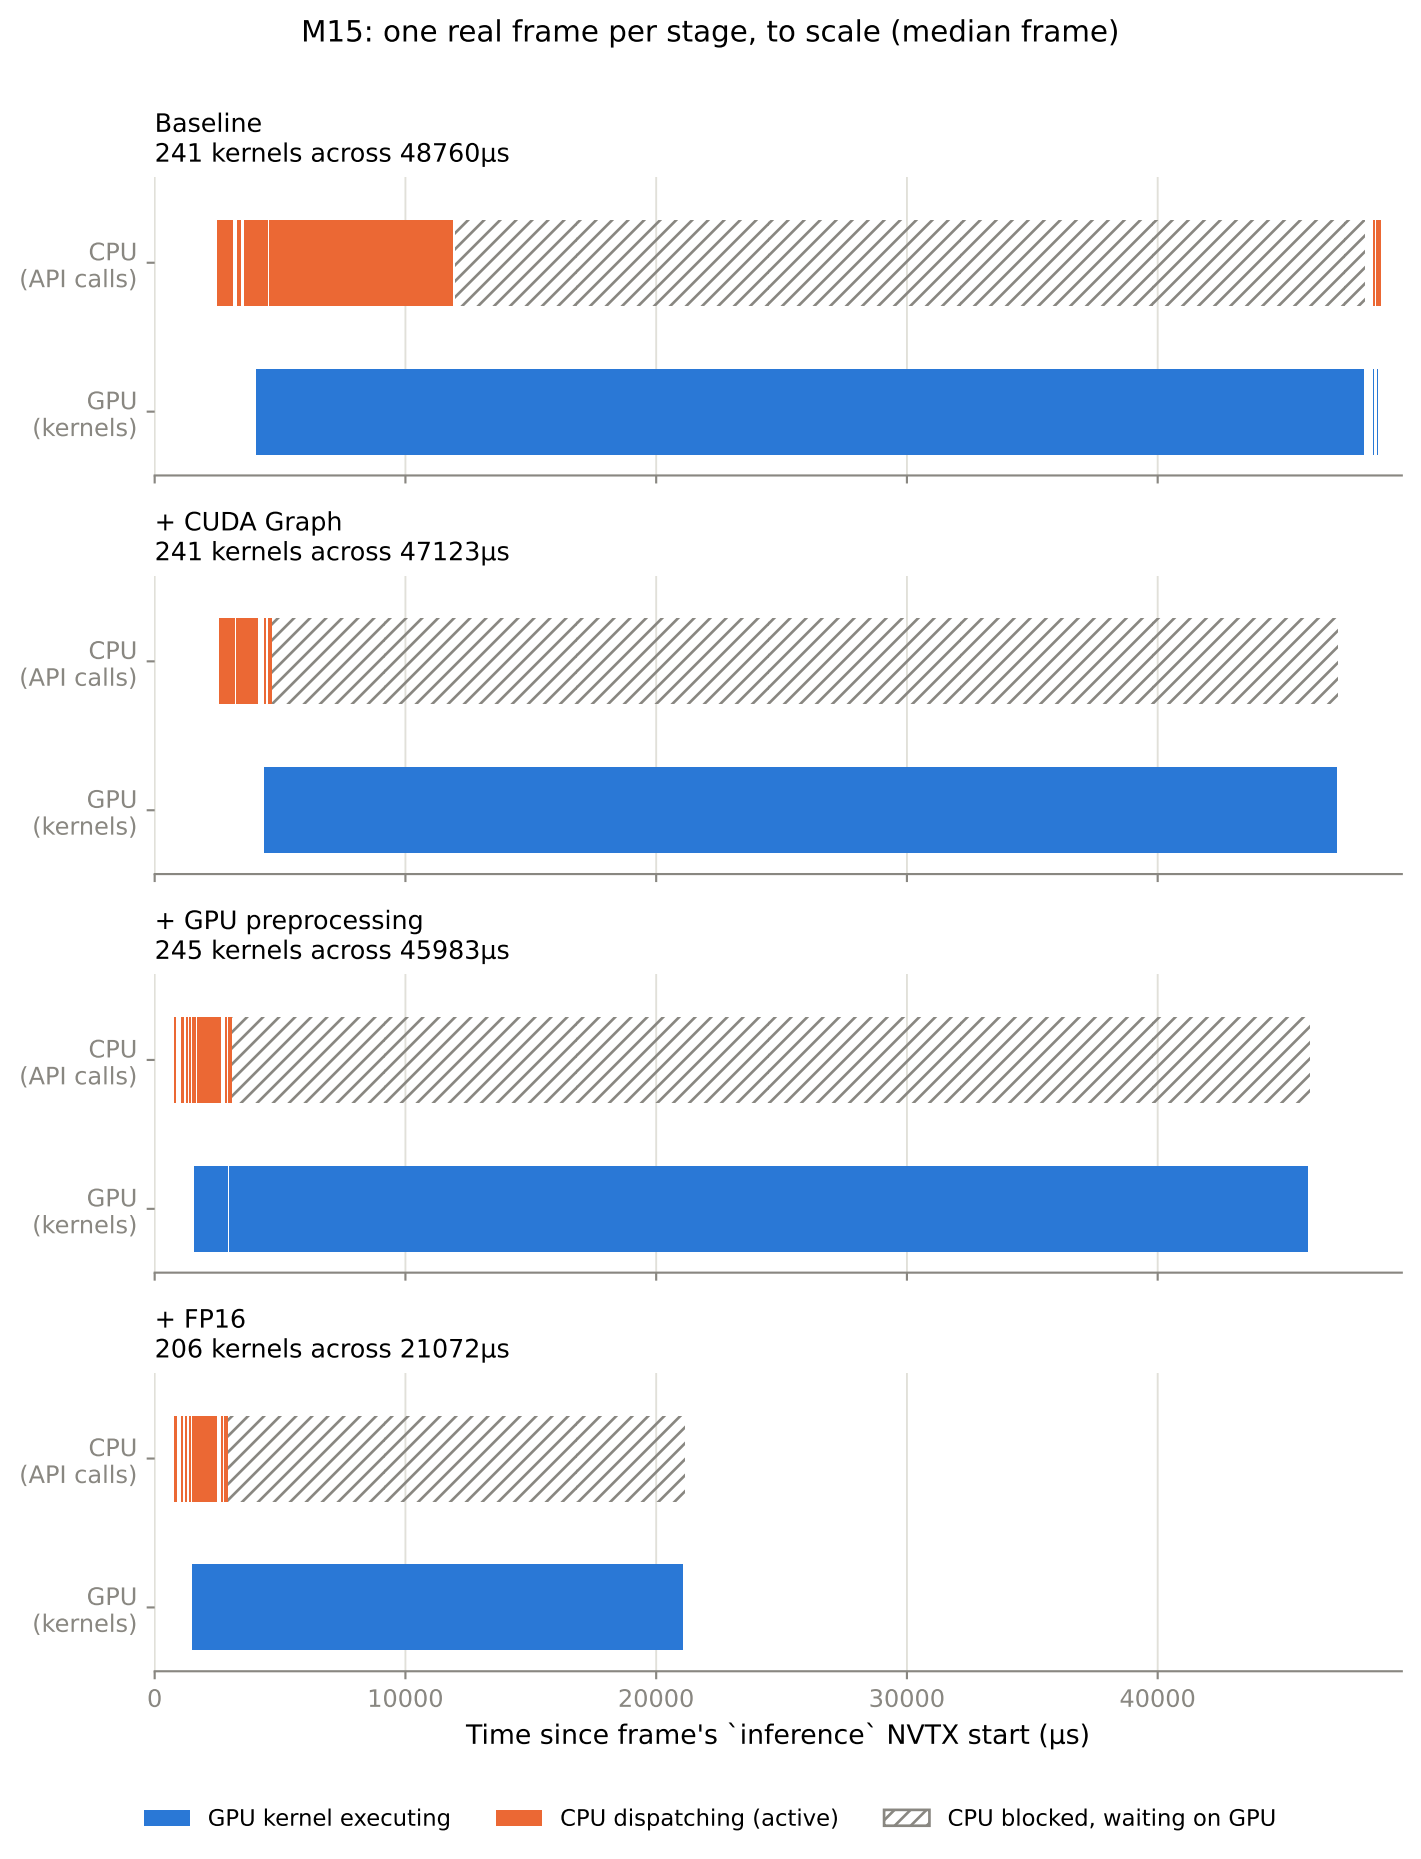
M15: one real frame per stage, to scale (median frame)
Baseline
241 kernels across 48760µs
CPU (API calls)
GPU (kernels)
+ CUDA Graph
241 kernels across 47123µs
CPU (API calls)
GPU (kernels)
+ GPU preprocessing
245 kernels across 45983µs
CPU (API calls)
GPU (kernels)
+ FP16
206 kernels across 21072µs
CPU (API calls)
GPU (kernels)
0
10000
20000
30000
40000
Time since frame's `inference` NVTX start (µs)
GPU kernel executing
CPU dispatching (active)
CPU blocked, waiting on GPU