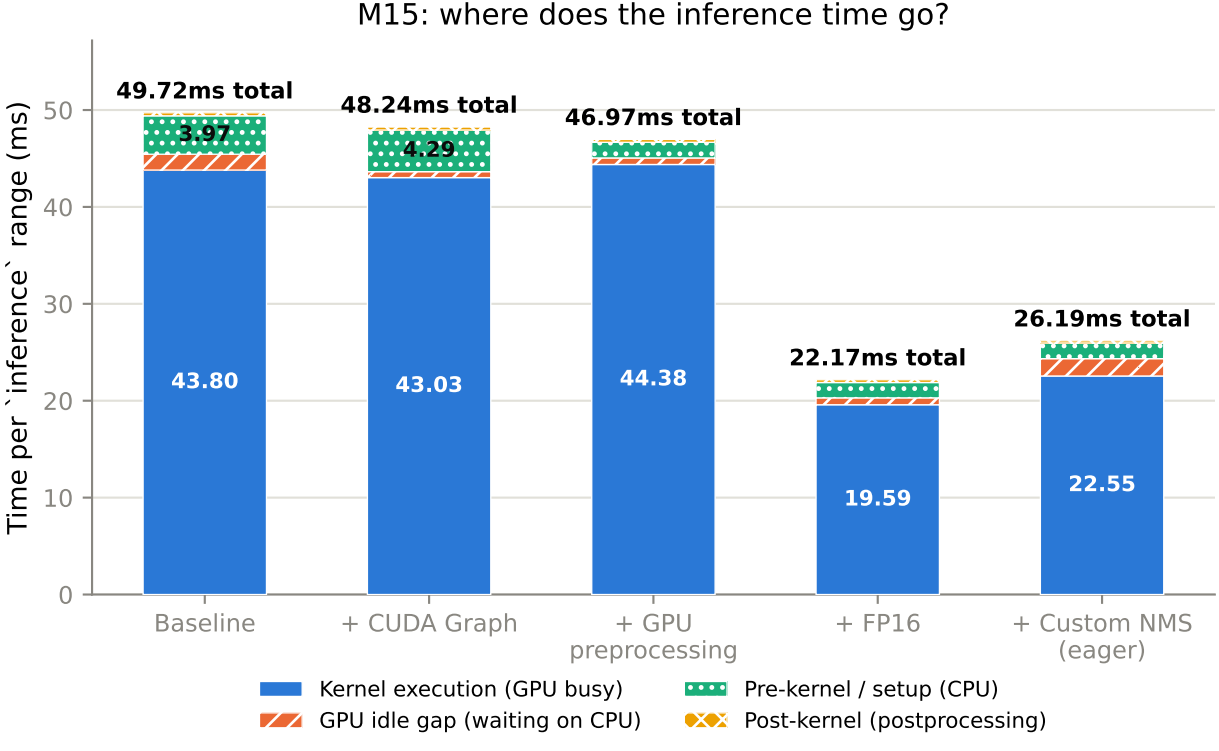
M15: where does the inference time go?
Time per `inference` range (ms)
0
10
20
30
40
50
49.72ms total
48.24ms total
46.97ms total
22.17ms total
26.19ms total
3.97
4.29
43.80
43.03
44.38
19.59
22.55
Baseline
+ CUDA Graph
+ GPU preprocessing
+ FP16
+ Custom NMS (eager)
Kernel execution (GPU busy)
GPU idle gap (waiting on CPU)
Pre-kernel / setup (CPU)
Post-kernel (postprocessing)