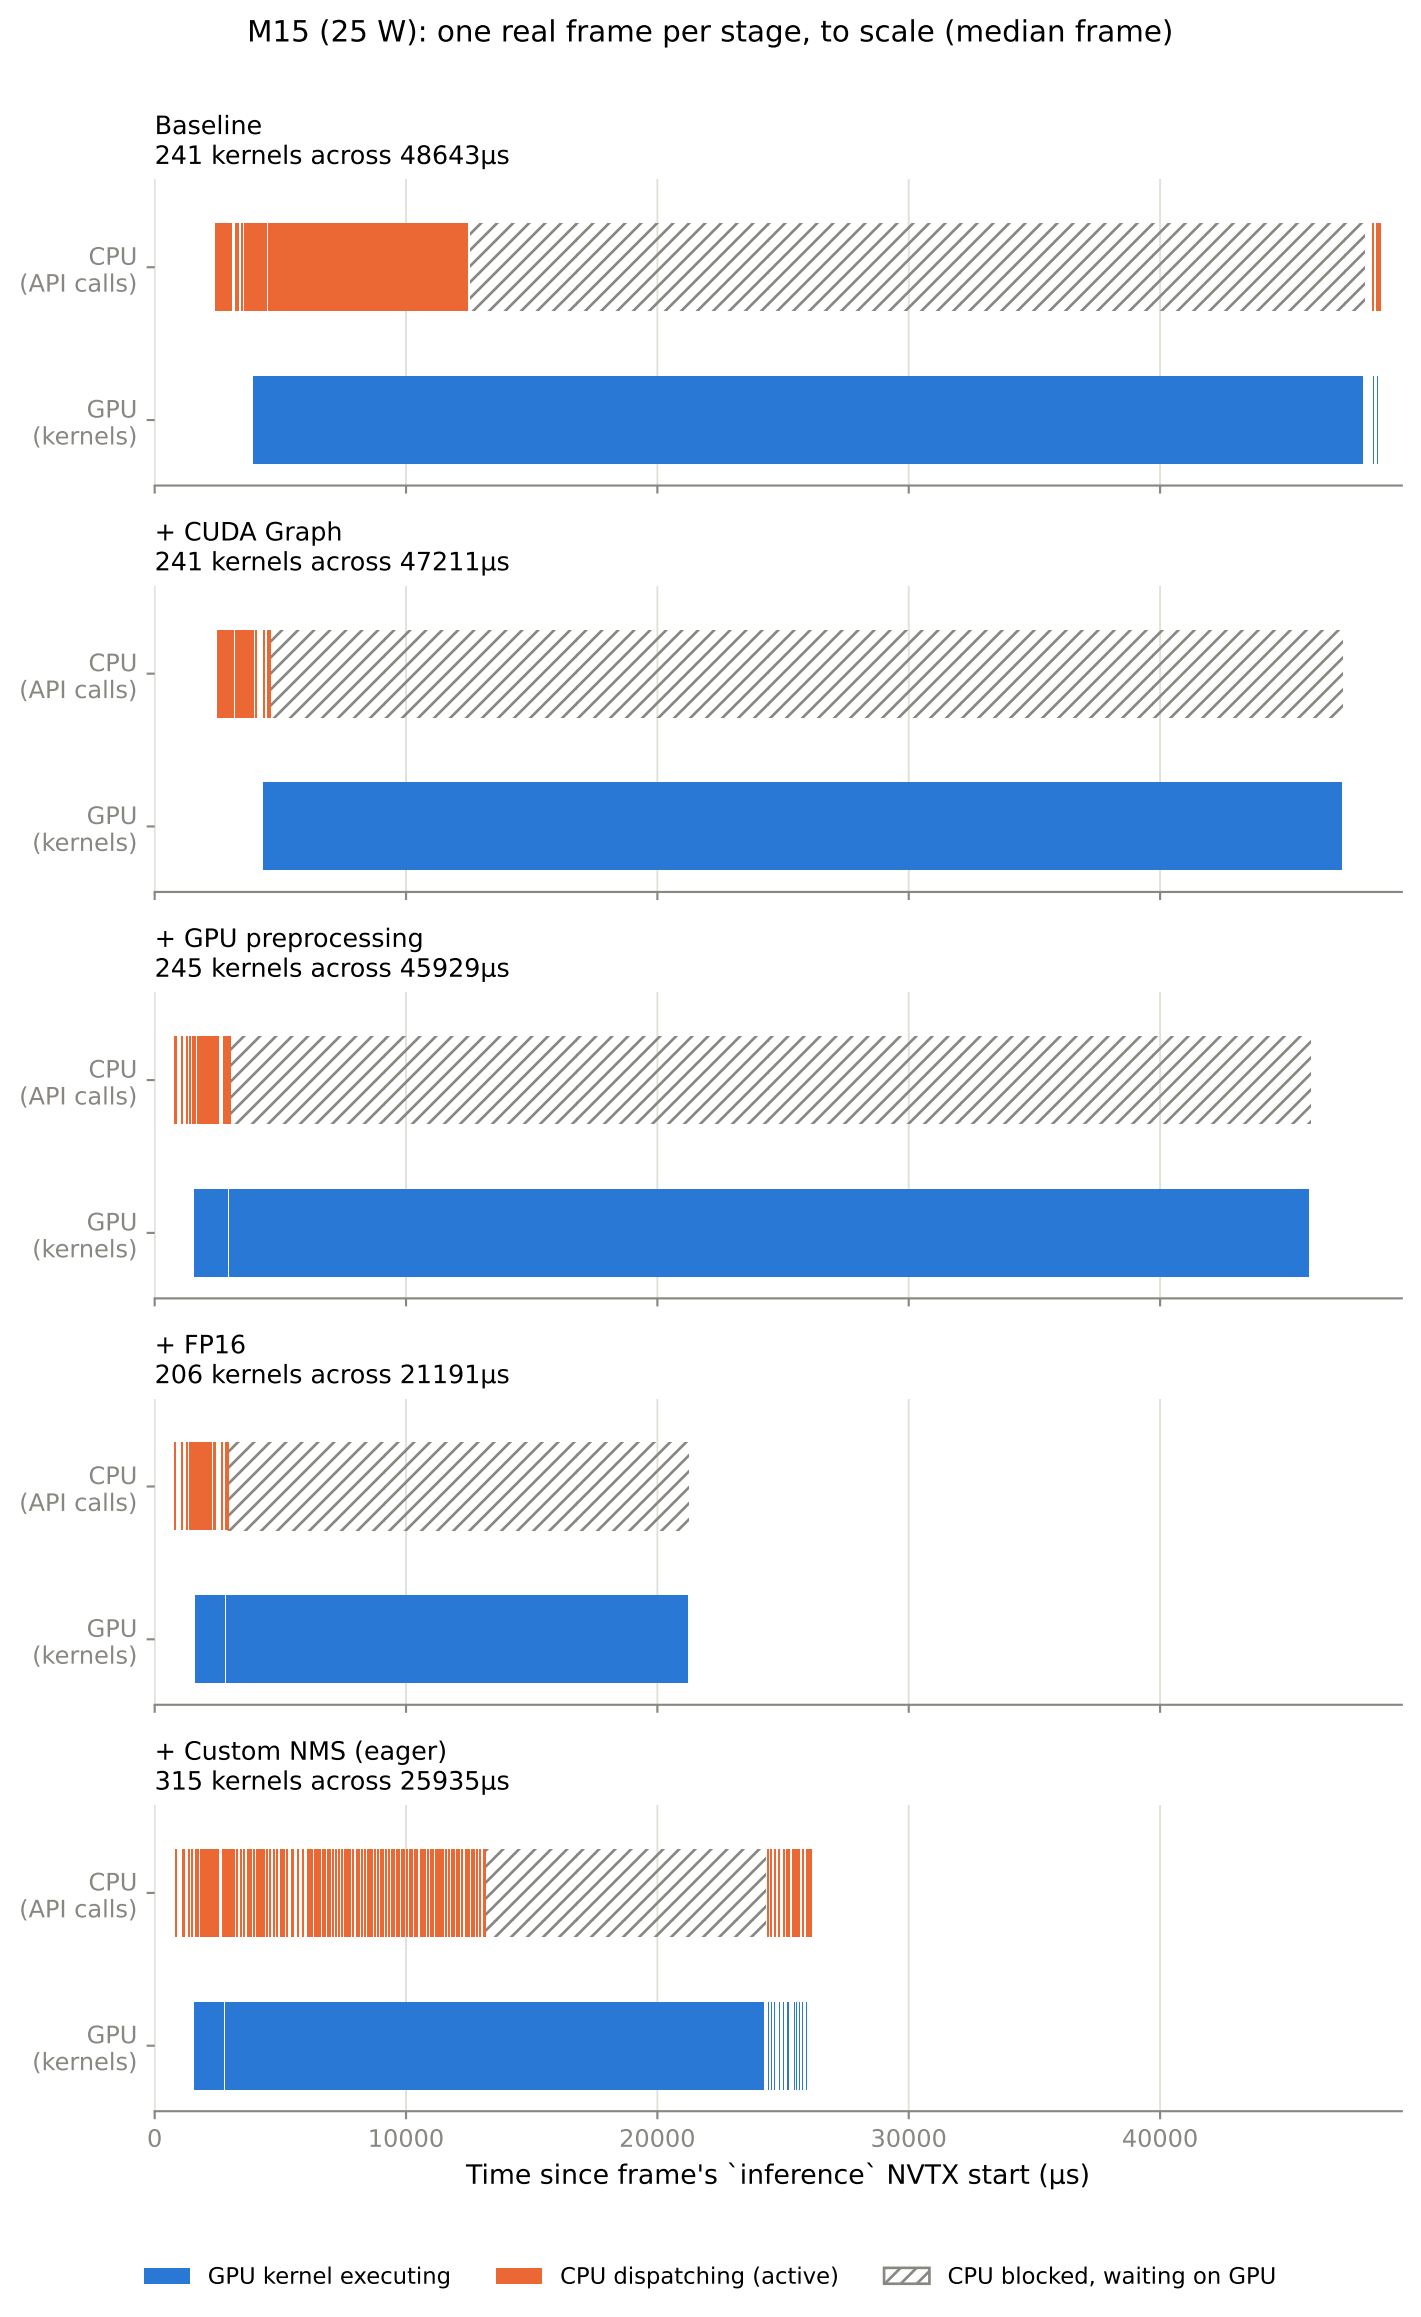
M15 (25 W): one real frame per stage, to scale (median frame)
Baseline
241 kernels across 48643µs
CPU (API calls)
GPU (kernels)
+ CUDA Graph
241 kernels across 47211µs
+ GPU preprocessing
245 kernels across 45929µs
+ FP16
206 kernels across 21191µs
+ Custom NMS (eager)
315 kernels across 25935µs
0
10000
20000
30000
40000
Time since frame's `inference` NVTX start (µs)
GPU kernel executing
CPU dispatching (active)
CPU blocked, waiting on GPU