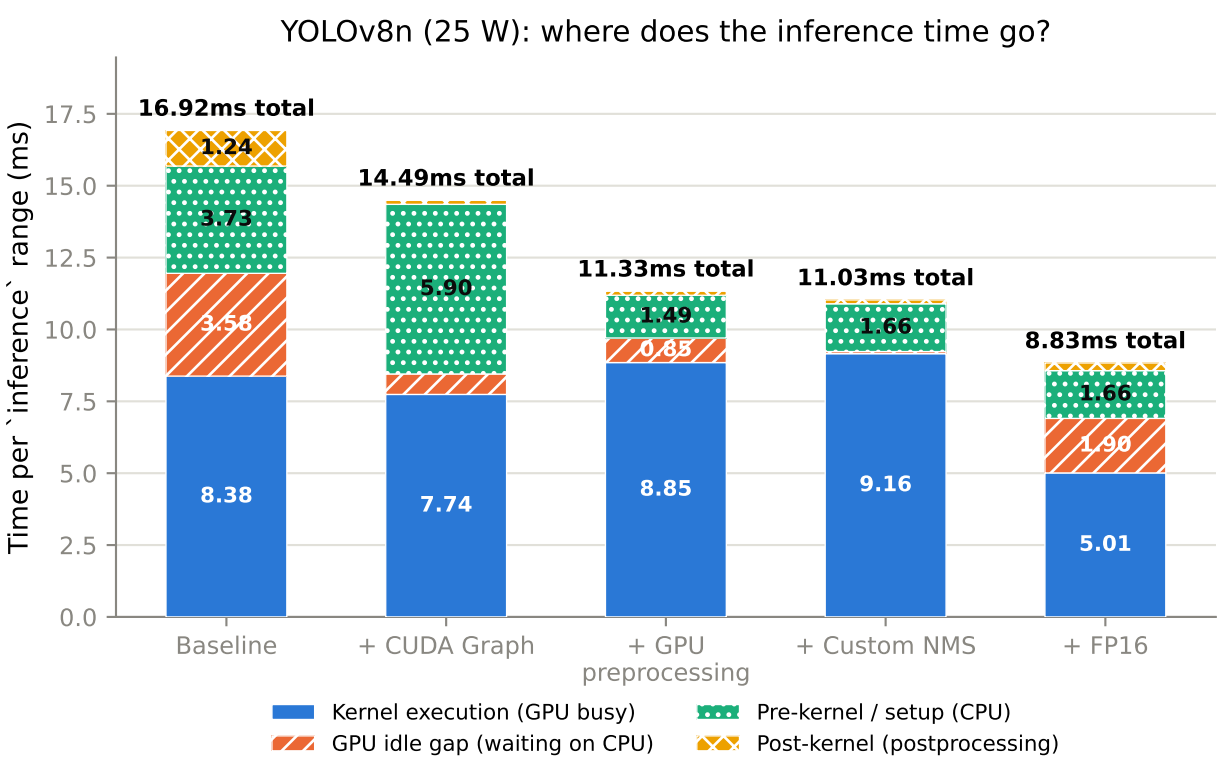
YOLOv8n (25 W): where does the inference time go?
Time per `inference` range (ms)
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
16.92ms total
1.24
3.73
3.58
8.38
14.49ms total
5.90
7.74
11.33ms total
1.49
0.85
8.85
11.03ms total
1.66
9.16
8.83ms total
1.66
1.90
5.01
Baseline
+ CUDA Graph
+ GPU preprocessing
+ Custom NMS
+ FP16
Kernel execution (GPU busy)
GPU idle gap (waiting on CPU)
Pre-kernel / setup (CPU)
Post-kernel (postprocessing)