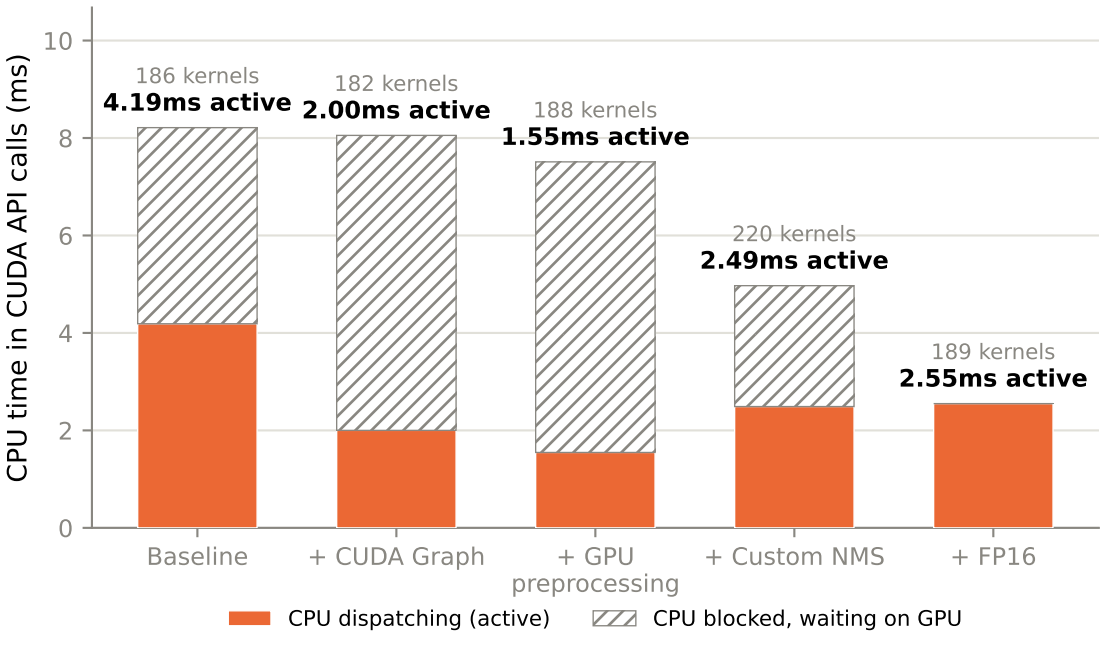
YOLOv8n (25 W): CPU time in CUDA API calls
CPU time in CUDA API calls (ms)
0
2
4
6
8
10
186 kernels
4.19ms active
182 kernels
2.00ms active
188 kernels
1.55ms active
220 kernels
2.49ms active
189 kernels
2.55ms active
Baseline
+ CUDA Graph
+ GPU preprocessing
+ Custom NMS
+ FP16
CPU dispatching (active)
CPU blocked, waiting on GPU