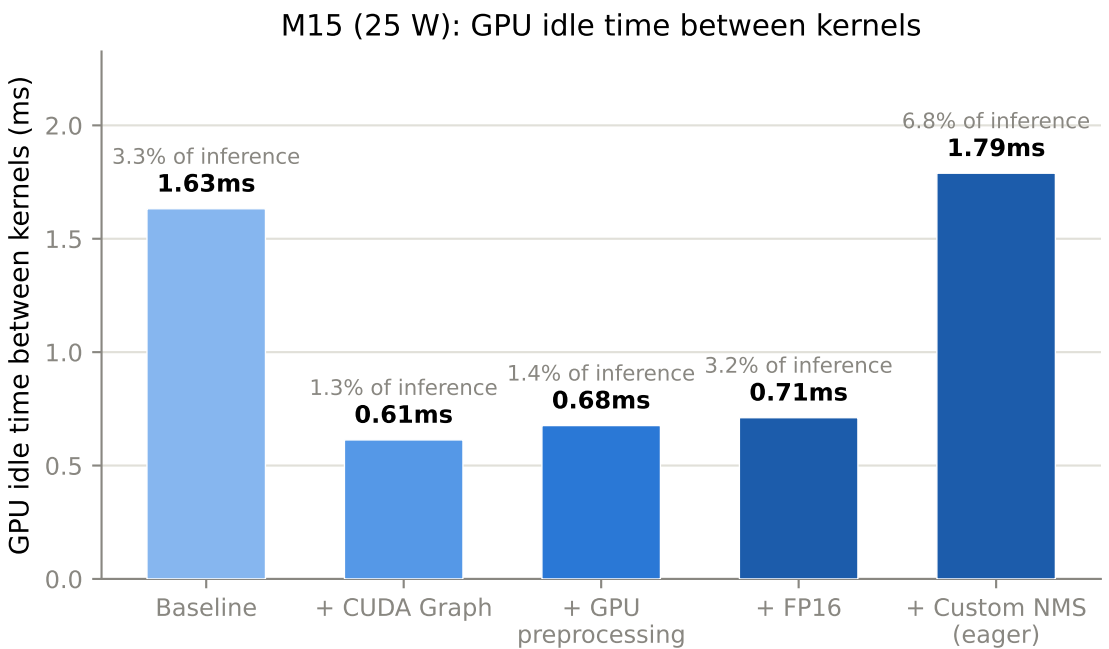
M15 (25 W): GPU idle time between kernels
GPU idle time between kernels (ms)
0.0
0.5
1.0
1.5
2.0
3.3% of inference
1.63ms
1.3% of inference
0.61ms
1.4% of inference
0.68ms
3.2% of inference
0.71ms
6.8% of inference
1.79ms
Baseline
+ CUDA Graph
+ GPU preprocessing
+ FP16
+ Custom NMS (eager)