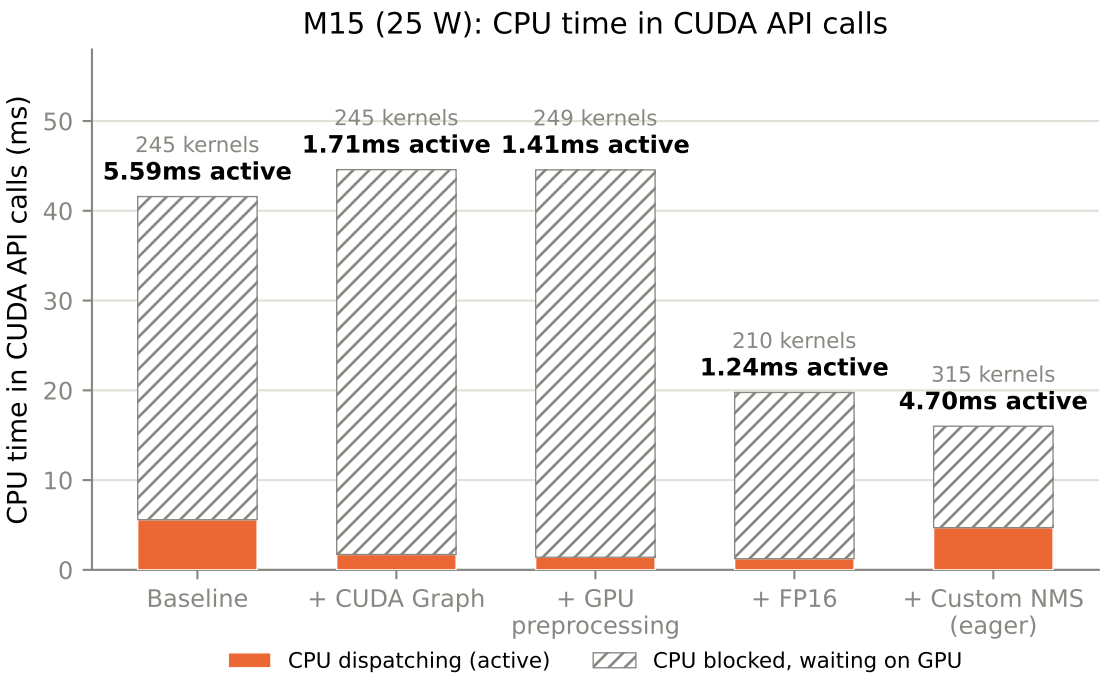
M15 (25 W): CPU time in CUDA API calls
CPU time in CUDA API calls (ms)
0
10
20
30
40
50
245 kernels
5.59ms active
245 kernels
1.71ms active
249 kernels
1.41ms active
210 kernels
1.24ms active
315 kernels
4.70ms active
Baseline
+ CUDA Graph
+ GPU preprocessing
+ FP16
+ Custom NMS (eager)
CPU dispatching (active)
CPU blocked, waiting on GPU