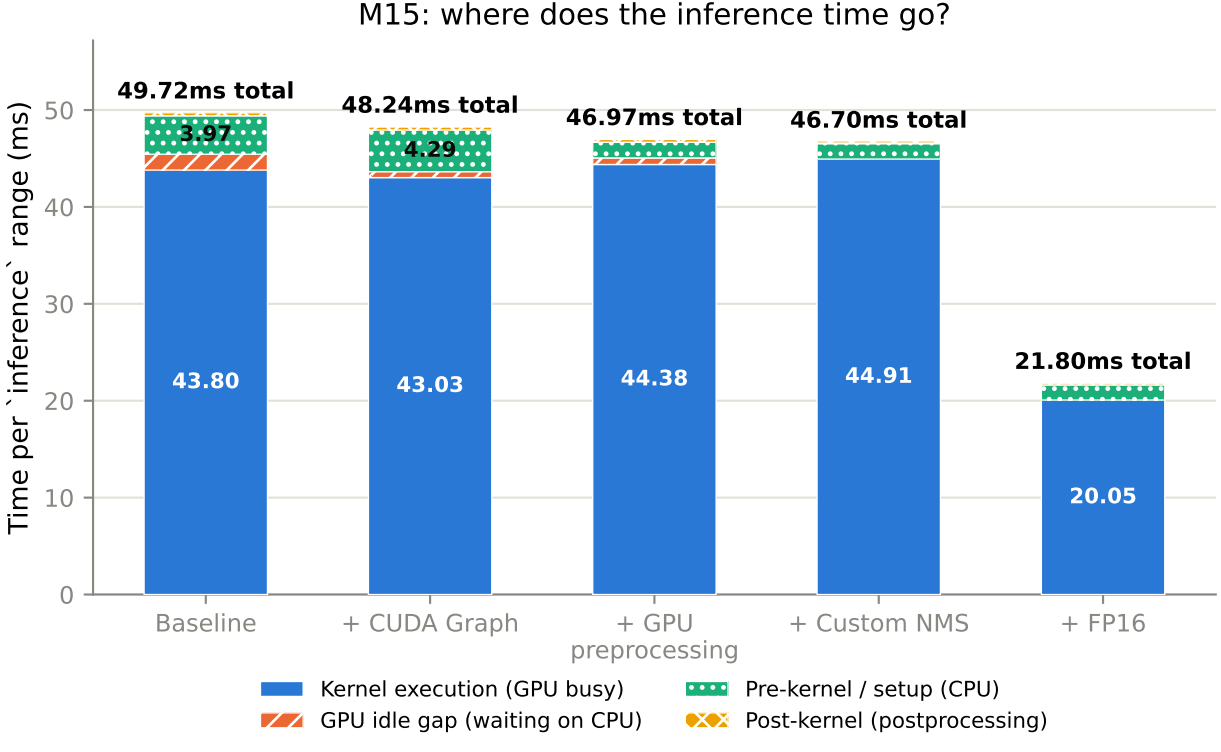
M15: where does the inference time go?
Time per `inference` range (ms)
0
10
20
30
40
50
49.72ms total
3.97
43.80
48.24ms total
4.29
43.03
46.97ms total
44.38
46.70ms total
44.91
21.80ms total
20.05
Baseline
+ CUDA Graph
+ GPU preprocessing
+ Custom NMS
+ FP16
Kernel execution (GPU busy)
GPU idle gap (waiting on CPU)
Pre-kernel / setup (CPU)
Post-kernel (postprocessing)