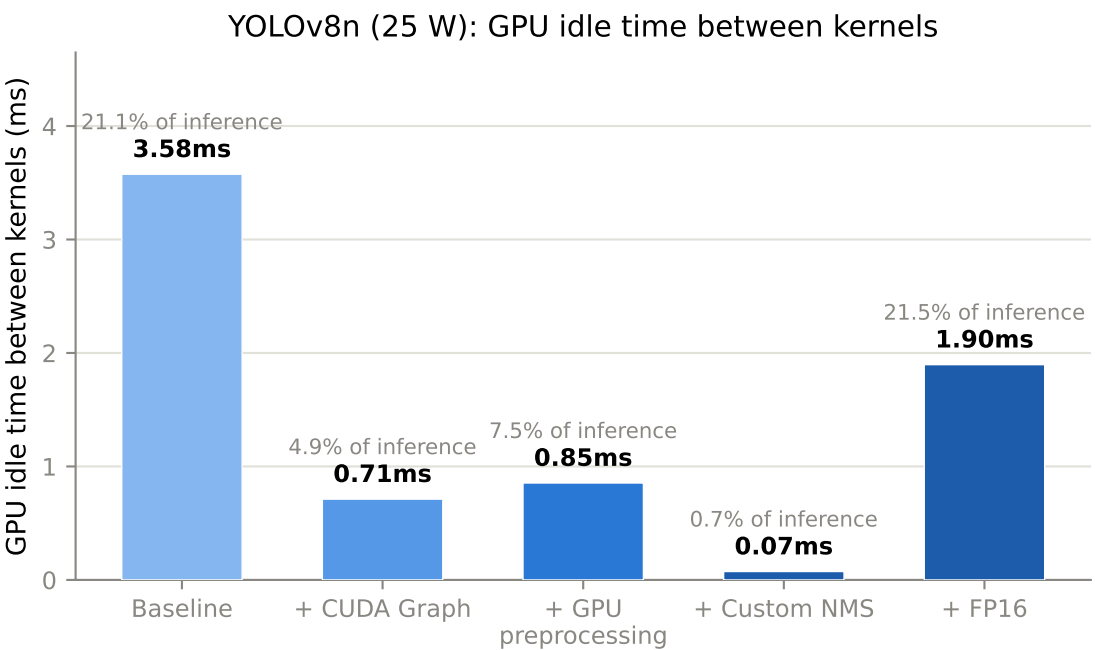
YOLOv8n (25 W): GPU idle time between kernels
GPU idle time between kernels (ms)
0
1
2
3
4
21.1% of inference
3.58ms
4.9% of inference
0.71ms
7.5% of inference
0.85ms
0.7% of inference
0.07ms
21.5% of inference
1.90ms
Baseline
+ CUDA Graph
+ GPU preprocessing
+ Custom NMS
+ FP16