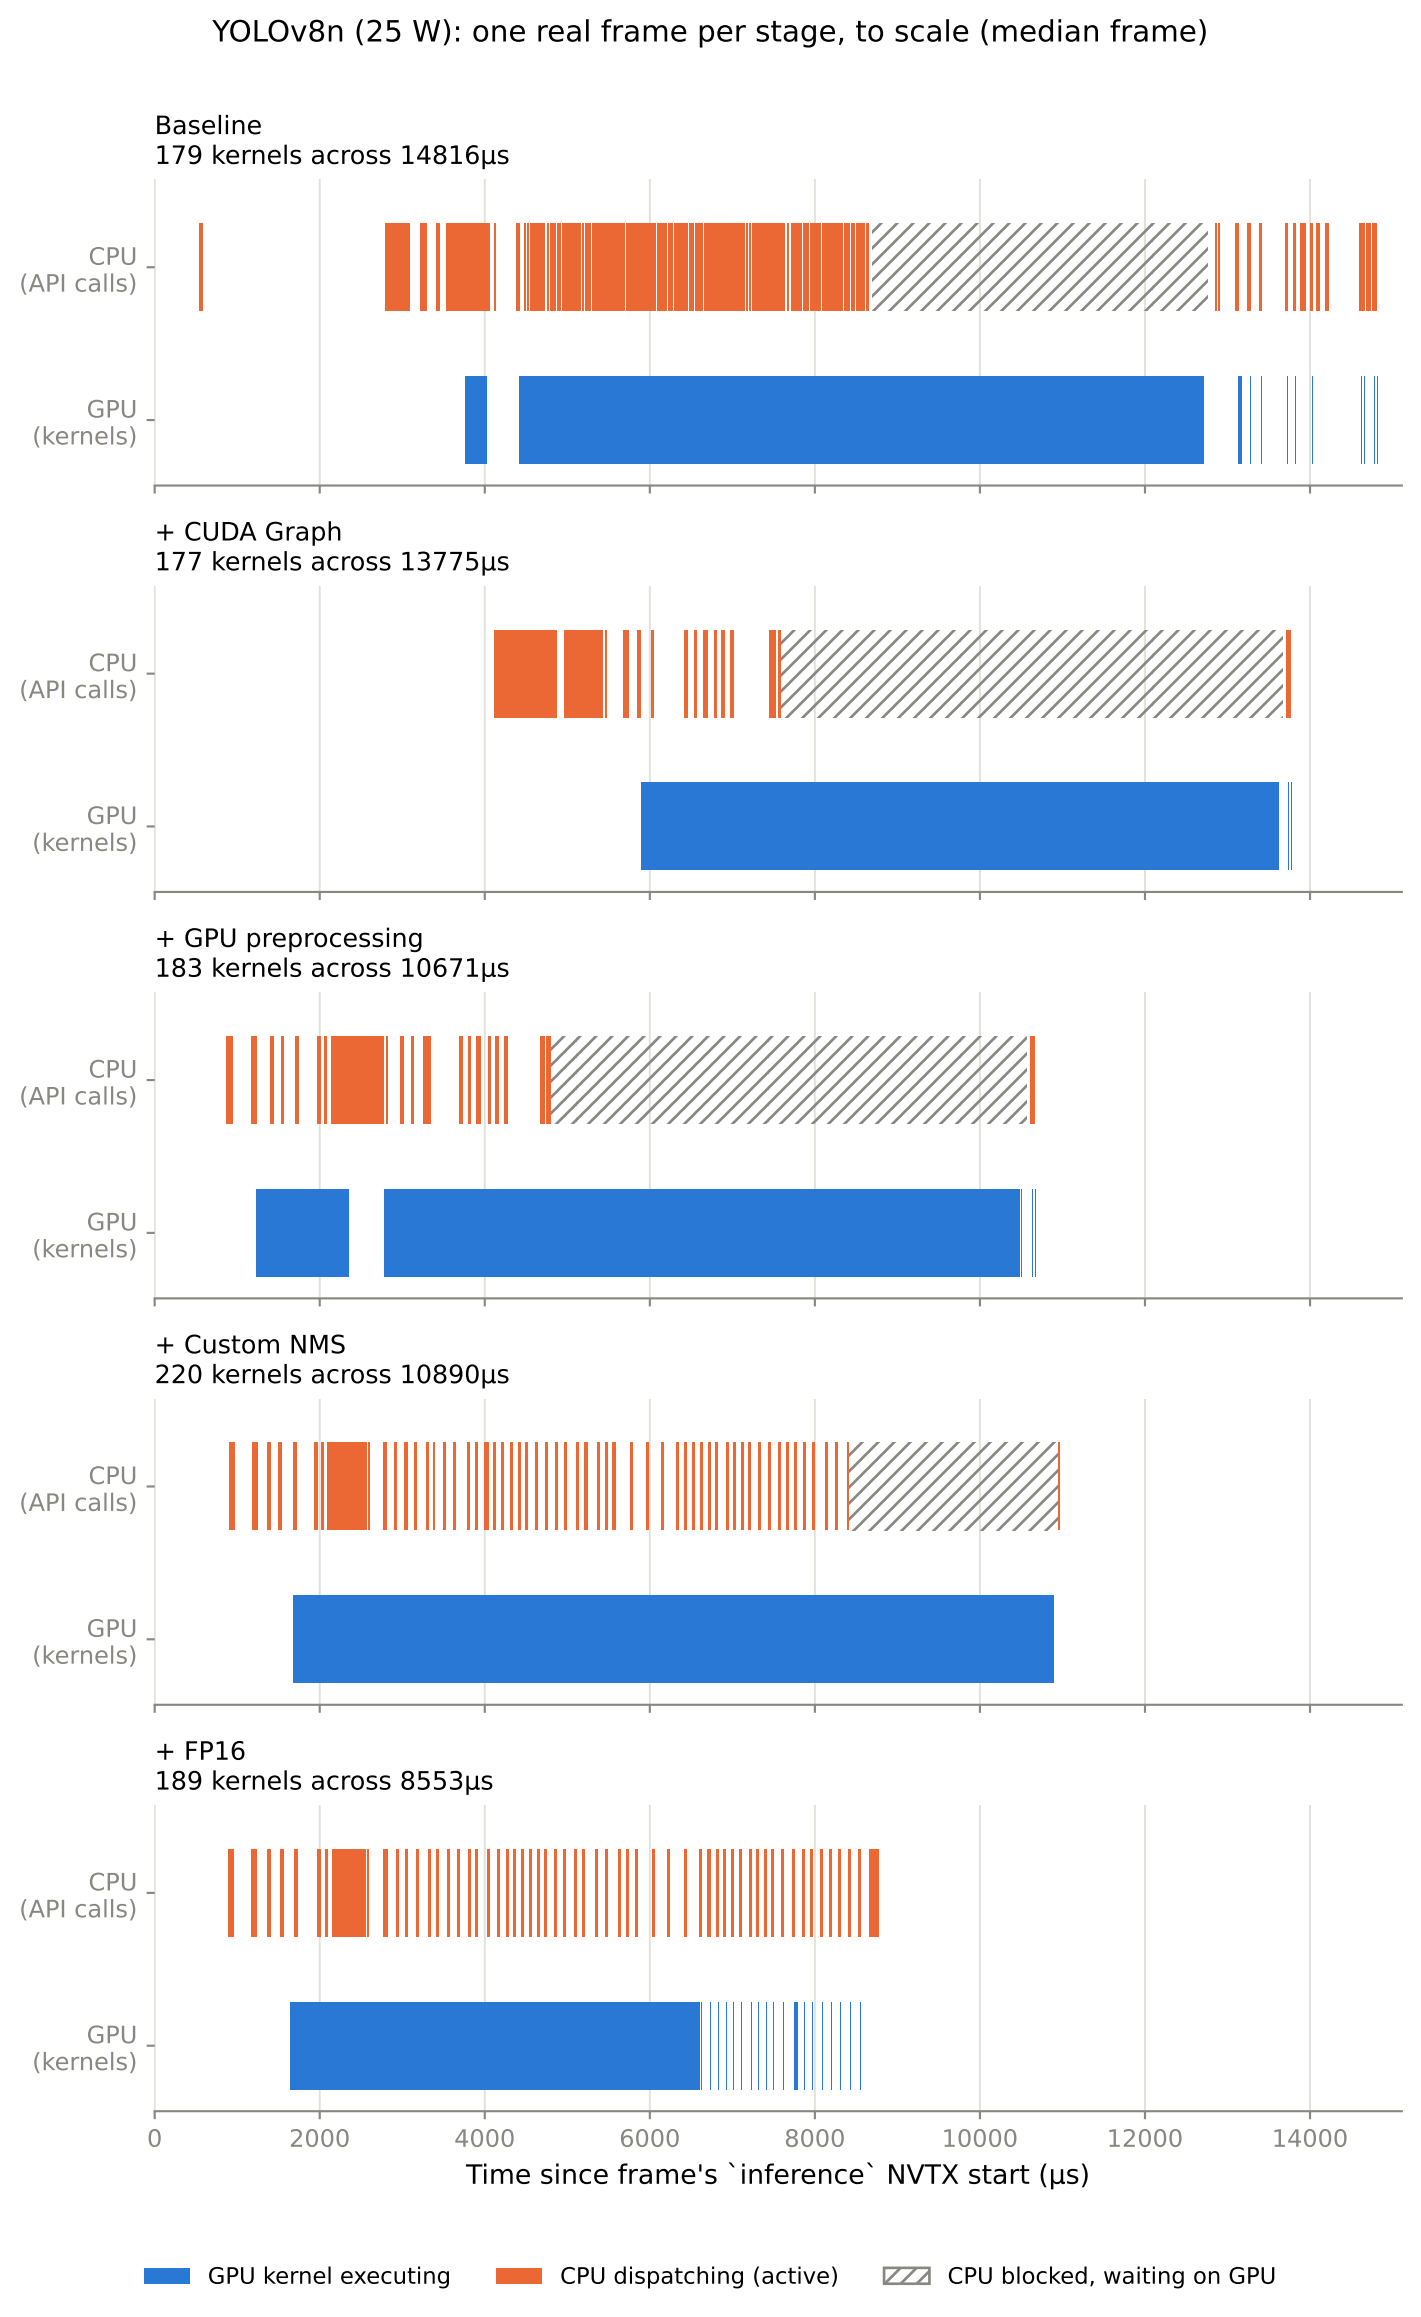
YOLOv8n (25 W): one real frame per stage, to scale (median frame)
Baseline
179 kernels across 14816µs
+ CUDA Graph
177 kernels across 13775µs
+ GPU preprocessing
183 kernels across 10671µs
+ Custom NMS
220 kernels across 10890µs
+ FP16
189 kernels across 8553µs
CPU (API calls)
GPU (kernels)
0
2000
4000
6000
8000
10000
12000
14000
Time since frame's `inference` NVTX start (µs)
GPU kernel executing
CPU dispatching (active)
CPU blocked, waiting on GPU