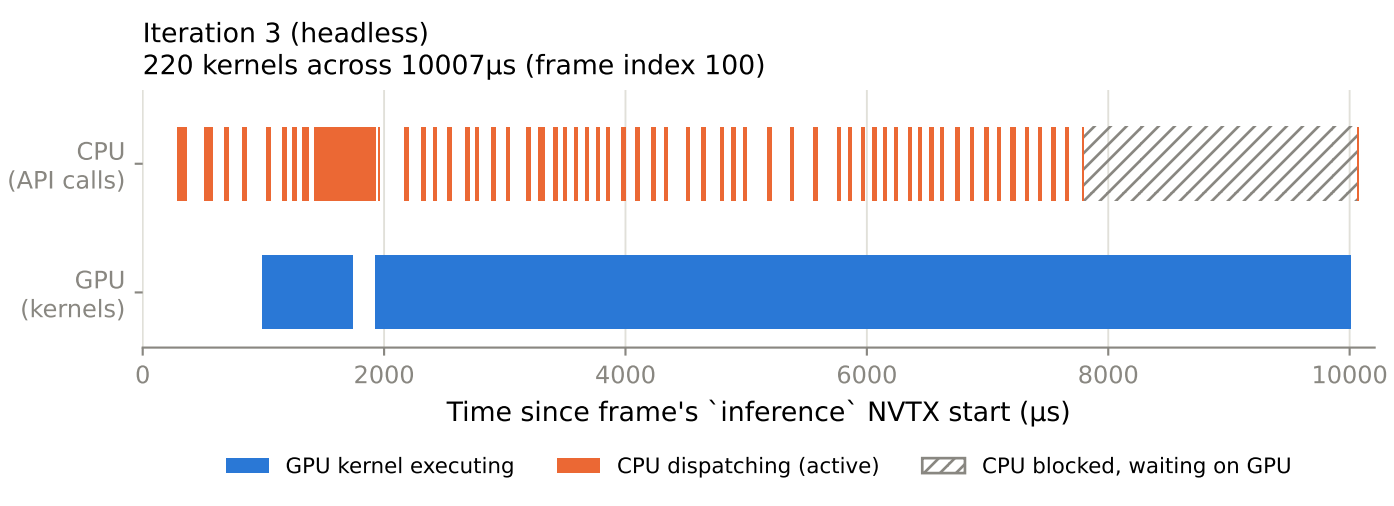
Iteration 3 (headless)
220 kernels across 10007µs (frame index 100)
CPU
(API calls)
GPU
(kernels)
0
2000
4000
6000
8000
10000
Time since frame's `inference` NVTX start (µs)
GPU kernel executing
CPU dispatching (active)
CPU blocked, waiting on GPU